

One Bad Apple Spoils the Bunch:

Exploiting P2P Applications to Trace and Profile Tor Users

Stevens Le Blond Pere Manils Abdelberi Chaabane
Mohamed Ali Kaafar Claude Castelluccia Arnaud Legout Walid Dabbous
I.N.R.I.A, France

Abstract

Tor is a popular low-latency anonymity network. However, Tor does not protect against the exploitation of an insecure application to reveal the IP address of, or trace, a TCP stream. In addition, because of the linkability of Tor streams sent together over a single circuit, tracing one stream sent over a circuit traces them all. Surprisingly, it is unknown whether this linkability allows in practice to trace a significant number of streams originating from secure (i.e., proxied) applications.

In this paper, we show that **linkability allows us to trace 193% of additional streams, including 27% of HTTP streams possibly originating from “secure” browsers.** In particular, we traced 9% of *all* Tor streams carried by our instrumented exit nodes. **Using BitTorrent as the insecure application, we design two attacks tracing BitTorrent users on Tor.** We run these attacks in the wild for 23 days and reveal 10,000 IP addresses of Tor users. Using these IP addresses, we then profile not only the BitTorrent downloads but also the websites visited *per country of origin of Tor users*. We show that BitTorrent users on Tor are over-represented in some countries as compared to BitTorrent users outside of Tor. By analyzing the type of content downloaded, we then explain the observed behaviors by the higher concentration of pornographic content downloaded at the scale of a country. Finally, we present results suggesting the existence of an underground BitTorrent ecosystem on Tor.

1 Introduction

Assume that a source wants to leak top secret documents anonymously. It is considered secure to do so through Tor using a privacy-enhancing browser plugin such as TorButton. However, **assume that, at the same time, this source uses another insecure application on Tor.** Is it then possible to associate the top secret documents with the IP address of the anonymous source? The answer to this question is yes!

By exploiting Tor’s design, one can indeed **exploit an insecure application to associate the usage of a secure application (e.g., the one leaking top secret documents) with the IP address of a Tor user.** This attack against Tor consists of two parts: (a) exploiting an insecure application to reveal the source IP address of, or *trace*, a

Tor user and (b) **exploiting Tor to associate the usage of a secure application with the IP address of a user** (revealed by the insecure application). As it is not a goal of Tor to protect against **application-level attacks**, Tor cannot be held responsible for the first part of this attack. However, because Tor’s design makes it possible to associate streams originating from secure application with traced users, the second part of this attack is indeed an attack against Tor. We call **the second part of this attack the bad apple attack**. (The name of this attack refers to the saying “one bad apple spoils the bunch.” We use this wording to illustrate that one insecure application on Tor may allow to trace other applications.)

This paper differs from the related work in three main aspects. First, we launched our attacks on the real Tor network for a substantial period of time and revealed 10,000 IP addresses of “anonymous” Tor users. To the best of our knowledge, it is the largest attack against the Tor network in number of revealed IP addresses.

Second, whereas **most attacks against Tor were targeted to Web browsers, we directly target P2P filesharing applications (i.e., BitTorrent) in this study.** BitTorrent traffic generates a significant fraction of Tor traffic in volume (more than 40%), making it a primary target for attackers. In addition, we show that 70% of BitTorrent users on Tor establish P2P connections *outside* of Tor, making most BitTorrent TCP connections (and traffic) invisible to the Tor network. Thus, the number of BitTorrent users on Tor is likely to be largely underestimated and so more Tor users are susceptible to our attacks.

Third, whereas the principle of the bad apple attack has been discussed in the past, it is an open question whether it allows to trace a significant number of streams originating from secure applications. Actually, that “many TCP streams can share one circuit” is listed as the fourth *improvement* of Tor over the old onion routing design [5] because it is supposed to “improve efficiency and anonymity.” We note that Roger Dingledine cited an initial version of this work [3, 10] to confirm the need to “brainstorm about ways to protect users even when their applications are handing over their sensitive information” on the website of the Tor Project [4]. The main contributions of this paper are as follow:

- We design two attacks against BitTorrent to reveal the IP address of BitTorrent users on Tor.
- We instrument six Tor exit nodes and launch our attacks on the real Tor network for a period of 23 days. We reveal 10,000 IP addresses of “anonymous” Tor users.
- We show that the bad apple attack allows us to trace 193% of additional streams as compared to BitTorrent streams, including 27% of HTTP streams. In total, we traced 9% of *all* Tor streams carried by our instrumented exit nodes.
- We profile BitTorrent and Web usage on Tor per country of origin (which would be impossible without first tracing Tor users). (a) We show that BitTorrent users on Tor are over-represented in some countries as compared to BitTorrent users that do not use Tor. (b) By analyzing the type of content downloaded, we explain this behavior by the higher concentration of pornographic content downloaded at the scale of a country. (c) We present results that suggest the existence of an underground BitTorrent ecosystem on Tor.

The rest of this paper is organized as follow. First, we briefly discuss the ethical and legal considerations of running attacks against production systems in Section 2. We then give an overview of Tor and explain how this design can lead to the bad apple attack when an insecure application reveals the source IP address of a Tor user in Section 3. In Section 4, we discuss our attack model, two attacks against BitTorrent, and how the bad apple attack specifically applies to BitTorrent. We show that the bad apple attack is severe enough to profile not only BitTorrent downloads but also the websites visited by Tor users in Section 5. We discuss the related work in Section 6. Finally, we summarize our contributions and give a few perspectives in Section 7.

2 Ethical and Legal Considerations

In order to comply with the legal and ethical aspects of privacy, we performed our analysis on the fly and do not store any nominative information such as IP addresses. We logged only the ASN and country of traced Tor users to be able to perform this study. In addition, we only present aggregated statistics as suggested by Loesing et al. in [9]. Finally, we have also been cautious not to inadvertently DoS Tor or BitTorrent infrastructures, or interfere with the normal usage of those systems.

3 Background

3.1 Tor

Tor is a low-latency anonymity network. As stated in the original paper, its main design goals are to prevent attackers from linking communication partners and from

linking multiple communications to or from a given user. Tor relies on an overlay network and on onion routing to anonymize TCP-based applications like web browsing and P2P filesharing. Tor explicitly made the design choice to support *only* TCP which “helped portability and deployability” [5].

When a client communicates with a server via Tor, she selects n nodes of the Tor system (where n is typically 3) and builds a circuit using those nodes. Messages are then encrypted n times, first with the key shared with the last node (called *exit node*) of the circuit, and subsequently with the shared keys of the intermediate nodes from $node_{n-1}$ to $node_1$. As a result, each intermediate node only knows its predecessor and successor, but no other nodes of the circuit. In addition, only the exit node is able to recover the original message.

To improve efficiency, Tor multiplexes several streams from the same source into a single circuit. Originally, onion routing used a separate circuit for each stream but it required multiple public key operations for every request [14]. It has also been argued that creating many circuits degraded privacy because it implied to contact more Tor nodes, some of which may be compromised. However, we will show that, when several streams are multiplexed into a single circuit, a single stream whose source IP address is revealed allows an attacker to associate many additional streams with the same traced user.

3.2 BitTorrent

BitTorrent is a popular Peer-to-peer (P2P) protocol for file replication. To download a content, a BitTorrent client first discovers peers sharing that content using centralized trackers, a distributed tracker (DHT), and peer exchange (PEX).

Trackers are servers storing content identifiers and for each identifier, a list of peers distributing the corresponding content. A peer subscribes its IP/port to the tracker for a given content identifier and requests a list of peers for that content when it starts downloading and then periodically after that (e.g., every 10 minutes). Communications with trackers is typically done in clear over TCP (i.e., through Tor) therefore a malicious exit node can tamper with the lists of peers returned by centralized trackers. (Furthermore, centralized tracking is sometimes also done over UDP with consequences similar to those we discuss after.)

In addition to centralized trackers, BitTorrent clients can also use a decentralized tracker based on a Distributed Hash Table (DHT tracking). Whereas DHT tracking is often used in combination to centralized tracking, it can also be used alone with Magnet Links. BitTorrent Mainline DHT tracking works as follow. First, a BitTorrent client (DHT tracker) picks an iden-

tifier that is coded on the same space as content identifiers. Then, a peer interested in downloading a content uses the content identifier to locate the corresponding DHT tracker after which, it subscribes its IP/port to that tracker and requests a list of peers (just as with a centralized tracker). Communication with DHT trackers is done over UDP therefore *a Tor user may subscribe her public IP/port to the DHT tracker.*

As a BitTorrent client discovers peers, it tries establishing TCP connections and if successful, sends an application handshake containing the content identifier and in the case of an extended handshake, the listening port number. That P2P connection is also used for content distribution. Whether a P2P connection is established through Tor has a tremendous impact on performance and can be configured by the user therefore *a Tor user may establish P2P connections using her public IP/port.*

Finally, after using centralized or DHT tracking, more peers can be discovered using Peer Exchange (PEX). With PEX, users typically exchange lists of peers they are connected to over established P2P connections therefore *a Tor user may subscribe her public IP/port to her PEX partners.*

4 Attacking BitTorrent Users on Tor

4.1 Attack Model

All our attacks require to control one exit node in order to trace its Tor users. From January 15 to February 7th 2010 (23 days), we instrument and monitor six Tor exit nodes spread throughout the world (two in Asia, two in Europe, and two in the U.S) and launch the attacks described after. The first attack, the hijacking tracker's responses, also requires to control a BitTorrent peer publicly connectable so it can accept incoming TCP connections and receive BitTorrent handshake messages. For practical reasons, we performed this hijacking attack on only one of our exit nodes.

Hijacking tracker's responses exploits the fact that centralized tracking is done through Tor and the list of peers can be tampered with by a malicious exit node. One condition for this attack to trace BitTorrent users is that they do content distribution outside of Tor. We will see in Section 4.2.1 that this is the case for 70% of the BitTorrent users on Tor. The second attack exploits the fact that DHT tracking uses UDP and so is done outside of Tor. We note that if DHT tracking was instead done over TCP through Tor, it would still be possible to perform an hijacking attack as with centralized tracking. We suspect that it is possible to perform similar statistical attacks with PEX and centralized UDP tracking, however, we did not exploit them in this study.




Figure 1: For each day, this histogram shows the proportion of BitTorrent peers who use Tor only to connect to the centralized tracker (*Tracker only*) or also to distribute content (*Tracker+content*). *All* is the average over all days. 72% of peers use Tor only to connect to the tracker.

4.2 Tracing BitTorrent Streams on Tor

4.2.1 Hijacking Tracker's Responses

Hijacking the tracker responses consists in inserting the IP/port of a peer controlled by the attacker (malicious peer) into the list of peers returned by the tracker so the targeted user connects to the malicious peer. When BitTorrent peers use Tor only to connect to the centralized tracker and not to distribute content, they will connect to the malicious peer directly, i.e., outside of Tor, allowing the attacker to trace them. But how can the attacker distinguish between a direct connection from one that is going through Tor?

This can be done by collecting the IP addresses of Tor exit nodes (which are public) to check whether an incoming connection at the malicious peer originates from one of these addresses. If it does, the targeted user is distributing content through Tor. Otherwise, he uses Tor only to connect to the tracker and the connection to the malicious peer is direct (the source IP address field of the IP datagrams contains the real IP address of the targeted user).

In Fig. 1, we observe that most BitTorrent users use Tor only to connect to the centralized tracker, making the hijacking of tracker's responses a simple yet efficient attack. One explanation for this behavior may be that users distribute content outside of Tor to not degrade performance. In addition, Piatek et al. have showed that naive spies use tracker subscriptions as evidences of copyright infringement [13]. This might be another reason why users are mainly concerned by anonymizing their tracker subscription.

Because most P2P connections are established outside of Tor, most BitTorrent streams (and traffic) are invisible to the Tor network. Thus, the number of BitTorrent users on Tor is likely to be largely underestimated and so more Tor users are susceptible to our attacks.




Figure 2: The distribution of the listening port number subscribed to the Mainline DHT is uniform (top plot) and most torrents have few peers (bottom plot), resulting in a small number of collisions among listening ports for peers in the same torrent. *The listening port number constitutes a good identifier of a peer within a torrent.*

4.2.2 Statistical Exploitation of DHT Tracking

The second attack exploits DHT tracking. Because DHT tracking is carried over UDP (which Tor does not support) a BitTorrent client fails to connect to the DHT through Tor and reverts to using its *public* interface to publish his (public) IP address and listening port into the DHT. Say the targeted user is downloading content1. The IP address and listening port of that user are then stored on the peer responsible of tracking content1 in the DHT. But how can an attacker in Tor locate the peer storing the public IP address of the targeted user? And then, how to distinguish the IP address of the targeted user from the other IP addresses downloading content1?

To find the information of a targeted user in the DHT, we **use the content identifier and listening port number contained in the BitTorrent subscription to the centralized tracker and extended handshake messages.** When one of our exit nodes receives one of these messages, it immediately locates the peer tracking that content identifier and collects all the IP/port couples that have subscribed for that identifier. All peers who have subscribed for that content identifier are candidates to be associated with the Tor user. We validate that the listening port is a good identifier of a peer within a torrent in Fig. 2. We then associate the IP address of the only peer with a matching listening port to the targeted user. If there is no such peer or that there is more than one, we consider that we have failed to trace the targeted user.

4.3 The Bad Apple Attack Applied to BitTorrent

Now that we have seen how to trace BitTorrent streams, we describe the specifics of the bad apple attack for BitTorrent users. The bad apple attack can reveal the source IP address of streams in the same circuit, or even dif-




Figure 3: Number of non-BitTorrent streams traced per destination port number. *The bad apple attack traces a significant number of HTTP (port 80) and HTTPS (port 443) streams.*

ferent circuits. With BitTorrent as with any other application, revealing the source IP address of streams in the same circuit is straightforward; **when the source IP address of a BitTorrent stream is revealed, all streams multiplexed into the same circuit are associated with the same traced user.**

To reveal the source IP address of streams in different circuits, we exploit two patterns in BitTorrent signalling traffic. The first pattern is the peer identifier which is essentially a random string of 20 bytes. Thus, when we observe the same peer identifier in the BitTorrent messages of different circuits, we consider that these circuits have the same source. One problem with this first pattern is that it does not work when communication between peers is encrypted. To alleviate this issue, we also consider communication to an IP/port freshly returned in a tracker response as a second pattern. In particular, if in a circuit we observe that a peer initiates communication with an IP/port contained in a tracker response from another circuit, we link the two circuits. We note that the linkage of BitTorrent streams in different circuits is particularly severe because, when an attacker traces a Tor user, he can potentially associate past or future circuits with that user, without need to reveal his IP address a second time.

We found that the bad apple attack applied to BitTorrent allowed us to trace 193% of additional streams as compared to BitTorrent streams, including 27% of HTTP streams possibly originating from “secure” browsers. We show the number of non-BitTorrent streams traced per destination port number in Fig. 3.

5 Profiling Tor Users

By analyzing the traffic relayed by our exit nodes, we evaluate that 19% of all streams on Tor are BitTorrent streams. (We remark that this percentage is much higher than in McCoy et al. [11], suggesting that the number of BitTorrent users on Tor has increased since 2008.) We

Rank	#	%	Over	Country
1	958	14	0.9	US
2	937	13	5.6	Japan
3	887	13	2.8	Germany
4	369	5	1.3	France
5	354	5	1.8	Poland
6	236	3	0.9	Italy
7	232	3	0.6	UK
8	231	3	-	China
9	203	3	0.7	Canada
10	200	2	1.4	Russia

Rank	#	Over	Country	AS
1	362	4.7	Germany	Deutsche Telekom (3320)
2	274	5.7	Japan	NTT (4713)
3	177	2	Malaysia	TM Net (4788)
4	142	1	Italy	Telecom Italia (3269)
5	135	1.1	France	Orange (3215)
6	133	1	US	AT&T (7132)
7	128	4.5	Germany	Hanse Net (13184)
8	113	-	China	China Net (4134)
9	109	1.4	Poland	TP Net (5617)
10	104	1.8	Austria	UPC (6830)

Table 1: Popularity and over-representation of BitTorrent users on Tor per country (left) and AS (right).

successfully trace 9% of *all* Tor streams. In this section, we use the resulting 10,000 traced IP addresses to profile the BitTorrent downloads and websites visited by Tor users per country of origin.

5.1 BitTorrent Profiling

We start by investigating whether BitTorrent utilization per country and AS on Tor is different relatively to BitTorrent utilization outside of Tor. We then analyze the content downloaded by BitTorrent users on Tor for a few interesting countries and investigate the existence of an underground BitTorrent ecosystem on Tor. Finally, we analyze the websites visited by Tor users per country of origin.

5.1.1 Utilization per Country of Origin

To compare BitTorrent utilization on Tor and outside of Tor, we need a representative sample of BitTorrent users for each of these two utilizations. We see in Fig. 4 that, after 10 days, most of the top10 countries and AS (in number of IP addresses of traced Tor users) have reached their final rank. Therefore, we have a reasonably representative sample of the utilization of BitTorrent *on Tor* in these countries and ASes.

As a representative sample of the utilization of BitTorrent *outside of Tor*, we then use a sample of 10,000,000 IP addresses collected on August 22nd 2009 on the PirateBay, the largest BitTorrent tracker at that time. Indeed, the PirateBay was an order of magnitude larger than the second largest BitTorrent tracker at the time of the measurement [16] therefore we argue that a daily sample from the PirateBay is reasonably representative of the global utilization of BitTorrent outside of Tor. We refer to Le Blond et al. [2] for a description of the measurement methodology to collect this data.

Table 1 shows the popularity of BitTorrent users on Tor per country (left) and AS (right). The over-representation (Over) for a given country (resp. AS) is the fraction of BitTorrent IP addresses on Tor in that country (resp. AS) divided by the fraction of IP addresses outside of Tor in the same country (resp. AS).

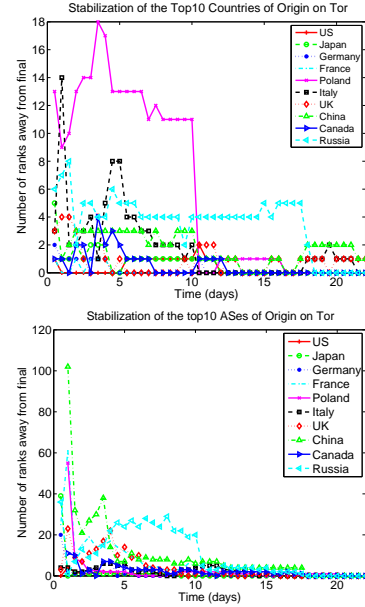


Figure 4: Stability of the top10 countries and ASes in cumulated number of IP addresses of traced Tor users. For each country and AS, we plot the evolution of the difference between the current rank and the rank after 23 days. *The duration of the measurement period is sufficient for the top10 countries and ASes to be reasonably representative of the overall utilization of BitTorrent on Tor.*

We do not show the over-representation in China because Chinese content are not generally tracked by the PirateBay, the tracker that we have used to capture the location of BitTorrent users outside of Tor. Hence, we greatly over-estimate the over-representation in China.

An over-representation of 0.9 in the US means that there is about the same fraction of US BitTorrent users on Tor as outside of Tor. And an over-representation of 5.6 in Japan means that there are 5.6 times more BitTorrent users from Japan on Tor than outside of Tor. In other words, whereas BitTorrent US users do not hide on Tor more than average, Japanese users strongly do. The




Figure 6: Distributions of BitTorrent content known in the *Public* and *Private* BitTorrent ecosystems, or *Unknown* from the rest of the world. A few content are unknown from the rest of the world.

reasons behind such behavior may be technological, political, sociological, etc.

Explaining Over-representations To investigate the over-representations observed in Section 5.1.1, we now analyze the types of content downloaded by Tor users from countries with very different over-representations. In particular, we had observed that US users were not over-represented on Tor whereas Japanese and German users were. In Fig. 5, we indeed see that US users are downloading a large variety of content as compared to Japanese users who mainly download Hentai (pornographic animes), and German users who mainly download pornographic movies. Therefore, we argue that the reasons for over-representations (at least in BitTorrent) are mainly sociological.

5.1.2 The Underground Ecosystem

BitTorrent comprises communities of users, or *ecosystems*, among which content are distributed [16], [2], [15]. There is one *public ecosystem* that can be accessed by all Internet users [2], [16] and several *private ecosystems* with many users whose access is restricted to registered users. Even though private ecosystems are more difficult to monitor than the public ecosystem, registration is generally open to everyone and a single registered user can in principle monitor the whole community [15]. Because even private ecosystems are indeed relatively easy to monitor, it is probable that some private ecosystems be known only by peculiar members (e.g., downloaders of child pornography) thus forming an *underground ecosystem*.

To investigate the existence of such an underground ecosystem on Tor, we check whether there is a subset of the content distributed on Tor that belongs neither to the public nor private ecosystems. We use .torrent files from 7,110,000 BitTorrent content: 6,800,000 from the public BitTorrent ecosystem [2, 16], and 310,000 from the private BitTorrent ecosystem [15]. We also search




Figure 7: Distributions of the categories of websites visited by Tor users per country of origin. *Tor* represents the overall distribution of requests for all Tor users (not necessarily traced). *BitTorrent* represents the overall distribution of requests for traced Tor users. The other bars represent the same information but per country of origin. *BitTorrent* users on Tor visit significantly more *Hacking* websites and significantly less *Search* and *Porn* than regular Tor users.

the missing .torrent files on Google. To the best of our knowledge, this is the largest collection of .torrent files ever assembled. We see in Fig. 6 that 3% of all the content distributed by BitTorrent users on Tor belong neither to the public nor private ecosystems. This result suggests the existence of an underground ecosystem on Tor. However, one would need to download these unknown content and to check them manually in order to determine whether they belong to regular private ecosystems for which we do not have the .torrent files or whether they are more sensitive.

5.2 Web Profiling

We have showed that the bad apple attack allowed us to trace 193% of additional streams as compared to BitTorrent streams, including 27% of HTTP streams possibly originating from “secure” (i.e., proxied) browsers. In the following, we use the web-filtering service of FortiGuard [7] to analyze the type of websites visited on Tor per country of origin.

We show the type of websites visited per country of origin in Fig. 7. Because they all are BitTorrent users, traced users differ from the average Tor users (*Tor*). In particular, around 50% of the requests are targeted to filesharing websites (*FileSharing*) such as *ThePirateBay*, *MegaDownload*, or *RapidShare*. Traced Tor users also visit significantly more hacking websites suggesting that they are interested in security. Finally, they also visit significantly less search and porn websites than average Tor users. This might be because they already rely on BitTorrent rather than the Web to search and download content, including pornographic material.



Figure 5: TagCloud of the US downloads (left), Japanese downloads (center), and German downloads (right). We extract the tags of the BitTorrent content in the public ecosystem and vary the font size to reflect the number of content whose tag matches those keywords. We increase the size of the keywords linearly with the frequency that they appear in the tags. *Japanese and German users use Tor to download much more pornographic material than US and other users.*

6 Related Work

Tor’s efficiency has required several adaptations to the original design of onion routing [14] whose impact on privacy are not well understood in practice [5]. In particular, we showed that the multiplexing of several streams into a single circuit can significantly degrade privacy.

To date, Tor measurement studies and attacks have been carried out in isolation. Measurement studies have documented *who is using Tor* and *how Tor is used* but without the ability of associating the two information, e.g., to profile Tor users [11]. Attacks have documented methodologies to associate the two information but without actually profiling Tor users [1, 6, 12]. This paper strikes a balance between the two by developing new attacks targeted to P2P applications, launching these attacks at a reasonable scale against the Tor network, and profiling Tor users.

6.1 Web-level Tracing Attacks

For simplicity, we assume that the attacker always controls an exit node in the attacks described hereafter.

We now describe attacks targeting Web applications to reveal the IP address of Tor users. FortConsult designed two attacks based on active content injection by the exit nodes to trace Tor users [6]. The first attacks consisted in Flash injection so that the targeted user connects to a server controlled by the attacker *outside* of Tor, hence exposing his IP address. (A cookie is used to associate that IP address to the stream in Tor.) The second attack consisted in injecting JavaScript to send the (local) IP address of the user over Tor. This study reported that whereas the first attack was effective, the second was not mainly because local IP addresses sent over Tor were not routable, e.g., 192.168.0.1.

Abbott et al. relied on JavaScript and HTML meta refresh tag to inject timing patterns [1]. The assumption being that users will leave that page open long enough so that pattern can be spotted by an entry node, also controlled by the attacker (thus tracing the user). It is interesting to note that even users having disabled active content are susceptible to HTML meta refresh tag injection.

To the best of our knowledge, we are the first to design attacks against P2P applications on Tor, to validate these attacks at a reasonable scale, and to demonstrate that one can associate many streams possibly originating from “secure” (i.e., proxied) browsers, with traced users. We remark that even though we have targeted P2P applications in this paper, the bad apple attack can originate from any insecure application.

6.2 Measurements Studies

The main measurement study of Tor that we are aware of has been made by McCoy et al. [11]. The authors provided interesting insights into *who is using Tor* and *how Tor is used and mis-used*. In particular, they showed that BitTorrent generates 40% of all traffic on Tor and claimed that it represents only 3% of all streams.

We complemented this measurement study in two important aspects. First, we showed that 70% of BitTorrent users establish P2P connections *outside* of Tor thus making most BitTorrent streams (and traffic) invisible to the Tor network. We argue that this finding makes of BitTorrent users a target of choice on Tor. Second, we launched attacks to profile Tor users. This profiling brought elements of answer to one important question raised by McCoy et al., that is, “... why there is such a large scale adoption of Tor in [...] specific countries, relative to Tor usage in other countries.” We showed that the answer to that question (at least for BitTorrent) is unlikely to be technological or political but in fact sociological.

7 Summary and Perspectives

Using BitTorrent as an insecure application, we designed two attacks, one consisting in hijacking tracker responses and one exploiting the statistical properties of the DHT, to trace BitTorrent streams on Tor. We then showed that the bad apple attack allows us to trace non-BitTorrent streams. In particular, we traced 193% of additional streams as compared to BitTorrent streams, including 27% of HTTP streams possibly originating from “secure” (i.e., proxied) browsers. In total, we traced 9% of *all* Tor streams carried by our instrumented exit nodes.

We ran these attacks in the wild for 23 days and reveal 10,000 IP addresses of Tor users. Using these IP addresses, we then profiled not only the BitTorrent downloads but also the websites visited per country of origin of Tor users. In particular, we showed that sociological reasons could explain the large number of Tor users in certain countries relatively to other. Finally, we presented results that suggest the existence of an underground BitTorrent ecosystem on Tor.

Defending against the bad apple attack is not straightforward. The most effective defense would be to have one stream per circuit, as in the original onion routing, however, performances issues make this defense unfeasible. Another defense would be to isolate streams by groups of destination port in different circuits, e.g., the secure and the insecure circuit. Destination ports known to be used by secure applications, e.g., 80, 22, would use the secure circuit thus limiting the risk that the source IP address of one stream in that circuit gets revealed by an insecure application. One weakness of this defense is that an attacker could trick an insecure application into connecting to a port that is usually used by a secure application, thus multiplexing the insecure stream into the secure circuit. Yet another defense would be to isolate each application into its own circuit, hence compartmenting the bad apple attack to the insecure application. However, modern operating systems lack a portable way to map an incoming stream to an application. We have discussed our results and possible solutions to address the bad apple attack with the Tor project.

We remark that even though the bad apple attack does not exist in application-level anonymity networks dedicated to a single application (e.g., OneSwarm [8]), the corpus of networking applications is too broad to practically build one network for each application. In this respect, we believe that we have validated an important attack against the design of modern anonymity networks and that we should defend against it to protect users privacy on the Internet.

Acknowledgements We warmly thank Keith Ross and Chao Zhang for sharing their identifiers of public and private content with us. We are also grateful to Roger Dingledine, Nick Mathewson, the anonymous reviewers, and our shepherd, Thorsten Holz, for their constructive comments.

References

- [1] T. Abbott, K. Lai, M. Lieberman, and E. Price. Browser-based Attacks on Tor. In *Proc. of PETS'07*, Ottawa, Canada, 2007.
- [2] S. L. Blond, A. Legout, F. Lefessant, W. Dabbous, and M. A. Kaafar. Spying the World from your Laptop - Identifying and Profiling Content Providers and Big Downloaders in BitTorrent. In *Proc. of LEET*, San Jose, CA, USA, 2010.
- [3] S. L. Blond, P. Manil, A. Chaabane, M. A. Kaafar, C. Castelluccia, A. Legout, and W. Dabbous. De-anonymizing BitTorrent Users on Tor. In *Proc. of NSDI'10, Poster session*, San Jose, CA, 2010.
- [4] R. Dingledine. Bittorrent Over Tor isn't a Good Idea. <https://blog.torproject.org/blog/bittorrent-over-tor-isnt-good-idea>, 2010.
- [5] R. Dingledine, N. Mathewson, and P. Syverson. Tor: the Second-generation Onion Router. In *Proc. of USENIX*, Boston, MA, 2004.
- [6] FortConsult. Practical Onion Hacking: Find the Real Address of Tor Clients. http://www.fortconsult.net/images/pdf/Practical_Onion_Hacking.pdf, 2006.
- [7] FortiGuard. Web filtering. <http://www.fortiguard.com/webfiltering/webfiltering.html>.
- [8] T. Isdal, M. Piatek, A. Krishnamurthy, and T. Anderson. Privacy-Preserving P2P Data Sharing With OneSwarm. In *Proc. of SIGCOMM*, Bangalore, India, 2010.
- [9] K. Loesing, S. Murdoch, and R. Dingledine. A Case Study on Measuring Statistical Data in the Tor Anonymity Network. In *Proc. of Financial Cryptography and Data Security '10*, 2010.
- [10] P. Manils, A. Chaabane, S. L. Blond, M. A. Kaafar, C. Castelluccia, A. Legout, and W. Dabbous. Compromising Tor Anonymity Exploiting P2P Information Leakage. Technical report, INRIA, 2010.
- [11] D. McCoy, T. Kohno, and D. Sicker. Shining Light in Dark Places: Understanding the Tor Network. In *Proc. of PETS'08*, Leuven, Belgium, 2008.
- [12] M. Perry. Securing the Tor Network. In *Proc. of Black Hat*, Las Vegas, NV, 2007.
- [13] M. Piatek, T. Kohno, and A. Krishnamurthy. Challenges and Directions for Monitoring P2P File Sharing Networks or Why My Printer Received a DMCA Takedown Notice. In *Proc. of HotSec*, San Jose, CA, 2008.
- [14] M. G. Reed, P. F. Syverson, and D. M. Goldschlag. Anonymous Connections and Onion Routing. *IEEE Journal on Selected Areas in Communications*, 16:482–494, 1998.
- [15] C. Zhang, P. Dhungel, Z. Liu, and K. W. Ross. BitTorrent Darknets. In *Proc. of INFOCOM*, San Jose, CA, USA, 2010.
- [16] C. Zhang, P. Dhungel, D. Wu, and K. W. Ross. Unraveling the BitTorrent Ecosystem. *TPDS*, 2010.